




Introduction to

AI-Driven HEP

S. A. Fard - School of Physics (IPM)






AI-Driven HEP 7



Session 7


Statistical Learning II



AI-Driven HEP 7

Reference

- [1] James, Gareth, et al. "Statistical learning." An introduction to statistical learning: With applications in Python. (2023)
- [2] Hastie, Trevor. "The elements of statistical learning: data mining, inference, and prediction." (2009)
- [3] Bishop, Christopher M., Pattern recognition and machine learning. springer, (2006)
- [4] Chollet, Francois. Deep learning with Python. (2021)
- [5] Shalev-Shwartz, Shai, and Shai Ben-David. Understanding machine learning: From theory to algorithms. (2014)
- [6] Tingjun Yang, A Brief Description of Data Products, ProtoDUNE Analysis Workshop, Jan 27, 2019, CERN
- [7] Gabriela Vitti Stenico, Development of the computing framework for monitoring the quality of ProtoDUNE off-line data, DUNE collaboration meeting (2025)




AI-Driven HEP 7

What does an Input look like

An array with shape (n, m, l, k)

(n, 1, 1, k)




(n, 1, l, k)



(n, m, l, k)







AI-Driven HEP 7

What does an Input look like

The organized (m, l, k) data format
no longer works






AI-Driven HEP 7

What does an Input look like

The organized (m, l, k) data format
no longer works



2.1 What Is Statistical Learning?


In order to motivate our study of statistical learning, we begin with a simple example. Suppose that we are statistical consultants hired by a client to investigate the association between advertising and sales of a particular product. The *Advertising* data set consists of the *sales* of that product in 200 different markets, along with advertising budgets for the product in each of those markets for three different media: *TV*, *radio*, and *newspaper*. The data are displayed in Figure 2.1. It is not possible for our client to directly increase sales of the product. On the other hand, they can control the advertising expenditure in each of the three media. Therefore, if we determine that there is an association between advertising and sales, then we can instruct our client to adjust advertising budgets, thereby indirectly increasing sales. In other words, our goal is to develop an accurate model that can be used to predict sales on the basis of the three media budgets.

In this setting, the advertising budgets are *input variables* while *sales* is an *output variable*. The input variables are typically denoted using the symbol X , with a subscript to distinguish them. So X_1 might be the *TV* budget, X_2 the *radio* budget, and X_3 the *newspaper* budget. The inputs go by different names, such as *predictors*, *independent variables*, *features*, or sometimes just *variables*. The output variable—in this case, *sales*—is often called the *response* or *dependent variable*, and is typically denoted using the symbol Y . Throughout this book, we will use all of these terms interchangeably.

More generally, suppose that we observe a quantitative response Y and p different predictors, $X_1, X_2, \dots, X_p$. We assume that there is some relationship between Y and $X = (X_1, X_2, \dots, X_p)$, which can be written in the very general form

$$Y = f(X) + \epsilon. \tag{2.1}$$


Set of words sequence



AI-Driven HEP 7

What does an Input look like

The organized (m, l, k) data format
no longer works



2.1 What Is Statistical Learning?


In order to motivate our study of statistical learning, we begin with a simple example. Suppose that we are statistical consultants hired by a client to investigate the association between advertising and sales of a particular product. The *Advertising* data set consists of the *sales* of that product in 200 different markets, along with advertising budgets for the product in each of those markets for three different media: *TV*, *radio*, and *newspaper*. The data are displayed in Figure 2.1. It is not possible for our client to directly increase sales of the product. On the other hand, they can control the advertising expenditure in each of the three media. Therefore, if we determine that there is an association between advertising and sales, then we can instruct our client to adjust advertising budgets, thereby indirectly increasing sales. In other words, our goal is to develop an accurate model that can be used to predict sales on the basis of the three media budgets.


In this setting, the advertising budgets are *input variables* while *sales* is an *output variable*. The input variables are typically denoted using the symbol X , with a subscript to distinguish them. So X_1 might be the *TV* budget, X_2 the *radio* budget, and X_3 the *newspaper* budget. The inputs go by different names, such as *predictors*, *independent variables*, *features*, or sometimes just *variables*. The output variable—in this case, *sales*—is often called the *response* or *dependent variable*, and is typically denoted using the symbol Y . Throughout this book, we will use all of these terms interchangeably.

More generally, suppose that we observe a quantitative response Y and p different predictors, $X_1, X_2, \dots, X_p$. We assume that there is some relationship between Y and $X = (X_1, X_2, \dots, X_p)$, which can be written in the very general form

$$Y = f(X) + \epsilon. \tag{2.1}$$


Set of words sequence






AI-Driven HEP 7

What does an Input look like




Graph format

	c1	c2	c3
r1	46	73	12
r2	89	19	35
r3	54	62	3




What does an Input look like




Graph format

	c1	c2	c3		
r1	46	73	12	11	$X = \begin{bmatrix} 46 \\ 73 \\ 12 \\ 89 \\ 19 \\ 35 \\ 54 \\ 62 \\ 3 \end{bmatrix}$
				12	
				13	
				21	
				22	
				23	
				31	
				32	
				33	
r2	89	19	35		
r3	54	62	3		



What does an Input look like



Graph format

	c1	c2	c3
r1	46	73	12
r2	89	19	35
r3	54	62	3

11
12
13
21
22
23
31
32
33


$X =$

46
73
12
89
19
35
54
62
3


111213212223313233

$A =$

110100000
111010000
011001000
100110100
010111010
001011001
000100110
000010111
000001011



What does an Input look like



Graph format

	c1	c2	c3
r1	46	73	12
r2	89	19	35
r3	54	62	3

11
12
13
21
22
23
31
32
33

$X =$

46
73
12
89
19
35
54
62
3


111213212223313233

$A =$

110100000
111010000
011001000
100110100
010111010
001011001
000100110
000010111
000001011


$E =$

(11,11) : w_1
(11,12) : w_2
⋮
⋮
(33,33) : w_{33}



What does an Input look like


Graph format



$$\begin{array}{l} 11 \\ 12 \\ 13 \\ 21 \\ 22 \\ 23 \\ 31 \\ 32 \\ 33 \end{array} X = \begin{bmatrix} 46 \\ 73 \\ 12 \\ 89 \\ 19 \\ 35 \\ 54 \\ 62 \\ 3 \end{bmatrix}$$

Relax the condition

$$\begin{array}{l} 11 \ 12 \ 13 \ 21 \ 22 \ 23 \ 31 \ 32 \ 33 \\ A = \begin{bmatrix} 1 & 0 & 0 & 1 & 0 & 1 & 0 & 0 & 0 \\ 1 & 1 & 1 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 1 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 1 & 0 & 1 & 1 \\ 0 & 1 & 1 & 0 & 1 & 1 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 & 0 & 0 & 1 & 1 & 0 \\ 0 & 1 & 0 & 1 & 1 & 0 & 0 & 1 & 1 \\ 0 & 0 & 0 & 1 & 0 & 1 & 1 & 1 & 1 \end{bmatrix} \end{array}$$



AI-Driven HEP 7

What does an Input look like

Text format

From natural languages to machine language

2.1 What Is Statistical Learning?

In order to motivate our study of statistical learning, we begin with a simple example. Suppose that we are statistical consultants hired by a client to investigate the association between advertising and sales of a particular product. The **Advertising** data set consists of the **sales** of that product in 200 different markets, along with advertising budgets for the product in each of those markets for three different media: **TV**, **radio**, and **newspaper**. The data are displayed in Figure 2.1. It is not possible for our client to directly increase sales of the product. On the other hand, they can control the advertising expenditure in each of the three media. Therefore, if we determine that there is an association between advertising and sales, then we can instruct our client to adjust advertising budgets, thereby indirectly increasing sales. In other words, our goal is to develop an accurate model that can be used to predict sales on the basis of the three media budgets.


In this setting, the advertising budgets are *input variables* while **sales** is an *output variable*. The input variables are typically denoted using the symbol X , with a subscript to distinguish them. So X_1 might be the **TV** budget, X_2 the **radio** budget, and X_3 the **newspaper** budget. The inputs go by different names, such as *predictors*, *independent variables*, *features*, or sometimes just *variables*. The output variable—in this case, **sales**—is often called the *response* or *dependent variable*, and is typically denoted using the symbol Y . Throughout this book, we will use all of these terms interchangeably.

More generally, suppose that we observe a quantitative response Y and p different predictors, $X_1, X_2, \dots, X_p$. We assume that there is some relationship between Y and $X = (X_1, X_2, \dots, X_p)$, which can be written in the very general form

$$Y = f(X) + \epsilon. \tag{2.1}$$

input
variable
output
variable

predictor
independent
variable
feature
variable
response
dependent
variable



What does an Input look like

Text format

From natural languages to machine language

2.1 What Is Statistical Learning?

In order to motivate our study of statistical learning, we begin with a simple example. Suppose that we are statistical consultants hired by a client to investigate the association between advertising and sales of a particular product. The **Advertising** data set consists of the **sales** of that product in 200 different markets, along with advertising budgets for the product in each of those markets for three different media: **TV**, **radio**, and **newspaper**. The data are displayed in Figure 2.1. It is not possible for our client to directly increase sales of the product. On the other hand, they can control the advertising expenditure in each of the three media. Therefore, if we determine that there is an association between advertising and sales, then we can instruct our client to adjust advertising budgets, thereby indirectly increasing sales. In other words, our goal is to develop an accurate model that can be used to predict sales on the basis of the three media budgets.

In this setting, the advertising budgets are *input variables* while **sales** is an *output variable*. The input variables are typically denoted using the symbol X , with a subscript to distinguish them. So X_1 might be the **TV** budget, X_2 the **radio** budget, and X_3 the **newspaper** budget. The inputs go by different names, such as *predictors*, *independent variables*, *features*, or sometimes just *variables*. The output variable—in this case, **sales**—is often called the *response* or *dependent variable*, and is typically denoted using the symbol Y . Throughout this book, we will use all of these terms interchangeably.


More generally, suppose that we observe a quantitative response Y and p different predictors, $X_1, X_2, \dots, X_p$. We assume that there is some relationship between Y and $X = (X_1, X_2, \dots, X_p)$, which can be written in the very general form

$$Y = f(X) + \epsilon. \tag{2.1}$$

input
variable
output
variable

predictor
independent
variable
feature
variable
response
dependent
variable

How much the sequence of the words affect our understanding ?



AI-Driven HEP 7

What does an Input look like

Text format

From natural languages to machine language

2.1 What Is Statistical Learning?

In order to motivate our study of statistical learning, we begin with a simple example. Suppose that we are statistical consultants hired by a client to investigate the association between advertising and sales of a particular product. The **Advertising** data set consists of the **sales** of that product in 200 different markets, along with advertising budgets for the product in each of those markets for three different media: **TV**, **radio**, and **newspaper**. The data are displayed in Figure 2.1. It is not possible for our client to directly increase sales of the product. On the other hand, they can control the advertising expenditure in each of the three media. Therefore, if we determine that there is an association between advertising and sales, then we can instruct our client to adjust advertising budgets, thereby indirectly increasing sales. In other words, our goal is to develop an accurate model that can be used to predict sales on the basis of the three media budgets.

In this setting, the advertising budgets are *input variables* while **sales** is an *output variable*. The input variables are typically denoted using the symbol X , with a subscript to distinguish them. So X_1 might be the **TV** budget, X_2 the **radio** budget, and X_3 the **newspaper** budget. The inputs go by different names, such as *predictors*, *independent variables*, *features*, or sometimes just *variables*. The output variable—in this case, **sales**—is often called the *response* or *dependent variable*, and is typically denoted using the symbol Y . Throughout this book, we will use all of these terms interchangeably.

More generally, suppose that we observe a quantitative response Y and p different predictors, $X_1, X_2, \dots, X_p$. We assume that there is some relationship between Y and $X = (X_1, X_2, \dots, X_p)$, which can be written in the very general form


$$Y = f(X) + \epsilon. \tag{2.1}$$

input
variable
output
variable

predictor
independent
variable
feature
variable
response
dependent
variable

How much the sequence of the words affect our understanding ?

How the similarities of the words are crucial ?



What does an Input look like

Text format

From natural languages to machine language

2.1 What Is Statistical Learning?

In order to motivate our study of statistical learning, we begin with a simple example. Suppose that we are statistical consultants hired by a client to investigate the association between advertising and sales of a particular product. The **Advertising** data set consists of the **sales** of that product in 200 different markets, along with advertising budgets for the product in each of those markets for three different media: **TV**, **radio**, and **newspaper**. The data are displayed in Figure 2.1. It is not possible for our client to directly increase sales of the product. On the other hand, they can control the advertising expenditure in each of the three media. Therefore, if we determine that there is an association between advertising and sales, then we can instruct our client to adjust advertising budgets, thereby indirectly increasing sales. In other words, our goal is to develop an accurate model that can be used to predict sales on the basis of the three media budgets.

In this setting, the advertising budgets are *input variables* while **sales** is an *output variable*. The input variables are typically denoted using the symbol X , with a subscript to distinguish them. So X_1 might be the **TV** budget, X_2 the **radio** budget, and X_3 the **newspaper** budget. The inputs go by different names, such as *predictors*, *independent variables*, *features*, or sometimes just *variables*. The output variable—in this case, **sales**—is often called the *response* or *dependent variable*, and is typically denoted using the symbol Y . Throughout this book, we will use all of these terms interchangeably.

More generally, suppose that we observe a quantitative response Y and p different predictors, $X_1, X_2, \dots, X_p$. We assume that there is some relationship between Y and $X = (X_1, X_2, \dots, X_p)$, which can be written in the very general form

$$Y = f(X) + \epsilon. \tag{2.1}$$


input
variable
output
variable

predictor
independent
variable
feature
variable
response
dependent
variable

How much the sequence of the words affect our understanding ?

How the similarities of the words are crucial ?

How the context of the words play roles ?



AI-Driven HEP 7

What does an Input look like

2.1 What Is Statistical Learning?


In order to motivate our study of statistical learning, we begin with a simple example. Suppose that we are statistical consultants hired by a client to investigate the association between advertising and sales of a particular product. The **Advertising** data set consists of the **sales** of that product in 200 different markets, along with advertising budgets for the product in each of those markets for three different media: TV, **radio**, and **newspaper**. The data are displayed in Figure 2.1. It is not possible for our client to directly increase sales of the product. On the other hand, they can control the advertising expenditure in each of the three media. Therefore, if we determine that there is an association between advertising and sales, then we can instruct our client to adjust advertising budgets, thereby indirectly increasing sales. In other words, our goal is to develop an accurate model that can be used to predict sales on the basis of the three media budgets.


In this setting, the advertising budgets are *input variables* while **sales** is an *output variable*. The input variables are typically denoted using the symbol X , with a subscript to distinguish them. So X_1 might be the **TV** budget, X_2 the **radio** budget, and X_3 the **newspaper** budget. The inputs go by different names, such as *predictors*, *independent variables*, *features*, or sometimes just *variables*. The output variable—in this case, **sales**—is often called the *response* or *dependent variable*, and is typically denoted using the symbol Y . Throughout this book, we will use all of these terms interchangeably.

More generally, suppose that we observe a quantitative response Y and p different predictors, $X_1, X_2, \dots, X_p$. We assume that there is some relationship between Y and $X = (X_1, X_2, \dots, X_p)$, which can be written in the very general form

$$Y = f(X) + \epsilon. \tag{2.1}$$

Text format





AI-Driven HEP 7

What does an Input look like

2.1 What Is Statistical Learning?


In order to motivate our study of statistical learning, we begin with a simple example. Suppose that we are statistical consultants hired by a client to investigate the association between advertising and sales of a particular product. The **Advertising** data set consists of the **sales** of that product in 200 different markets, along with advertising budgets for the product in each of those markets for three different media: TV, radio, and newspaper. The data are displayed in Figure 2.1. It is not possible for our client to directly increase sales of the product. On the other hand, they can control the advertising expenditure in each of the three media. Therefore, if we determine that there is an association between advertising and sales, then we can instruct our client to adjust advertising budgets, thereby indirectly increasing sales. In other words, our goal is to develop an accurate model that can be used to predict sales on the basis of the three media budgets.

In this setting, the advertising budgets are *input variables* while **sales** is an *output variable*. The input variables are typically denoted using the symbol X , with a subscript to distinguish them. So X_1 might be the **TV** budget, X_2 the **radio** budget, and X_3 the **newspaper** budget. The inputs go by different names, such as *predictors*, *independent variables*, *features*, or sometimes just *variables*. The output variable—in this case, **sales**—is often called the *response* or *dependent variable*, and is typically denoted using the symbol Y . Throughout this book, we will use all of these terms interchangeably.

More generally, suppose that we observe a quantitative response Y and p different predictors, $X_1, X_2, \dots, X_p$. We assume that there is some relationship between Y and $X = (X_1, X_2, \dots, X_p)$, which can be written in the very general form


$$Y = f(X) + \epsilon. \tag{2.1}$$

Text format



Order of total number of words in entire web $\sim 10^{15}$

Order of total number of unique English word $\sim 10^5$



AI-Driven HEP 7

What does an Input look like

2.1 What Is Statistical Learning?


In order to motivate our study of statistical learning, we begin with a simple example. Suppose that we are statistical consultants hired by a client to investigate the association between advertising and sales of a particular product. The *Advertising* data set consists of the *sales* of that product in 200 different markets, along with advertising budgets for the product in each of those markets for three different media: TV, radio, and newspaper. The data are displayed in Figure 2.1. It is not possible for our client to directly increase sales of the product. On the other hand, they can control the advertising expenditure in each of the three media. Therefore, if we determine that there is an association between advertising and sales, then we can instruct our client to adjust advertising budgets, thereby indirectly increasing sales. In other words, our goal is to develop an accurate model that can be used to predict sales on the basis of the three media budgets.

In this setting, the advertising budgets are *input variables* while *sales* is an *output variable*. The input variables are typically denoted using the symbol X , with a subscript to distinguish them. So X_1 might be the TV budget, X_2 the radio budget, and X_3 the newspaper budget. The inputs go by different names, such as *predictors*, *independent variables*, *features*, or sometimes just *variables*. The output variable—in this case, *sales*—is often called the *response* or *dependent variable*, and is typically denoted using the symbol Y . Throughout this book, we will use all of these terms interchangeably.

More generally, suppose that we observe a quantitative response Y and p different predictors, $X_1, X_2, \dots, X_p$. We assume that there is some relationship between Y and $X = (X_1, X_2, \dots, X_p)$, which can be written in the very general form


$$Y = f(X) + \epsilon. \tag{2.1}$$

Text format



Order of total number of words in entire web $\sim 10^{15}$

Order of total number of unique English word $\sim 10^5$




AI-Driven HEP 7

What does an Input look like

An array with shape (n, m, l, k)

Domain of the data




AI-Driven HEP 7

What does an Input look like

An array with shape (n, m, l, k)

Domain of the data

Incompleteness and generalisation




AI-Driven HEP 7


What does an Input look like

An array with shape (n, m, l, k)

Domain of the data

Incompleteness and generalisation






AI-Driven HEP 7


What does an Input look like

An array with shape (n, m, l, k)

Domain of the data

Incompleteness and generalisation







AI-Driven HEP 7

What does an Input look like


An array with shape (n, m, l, k)

Domain of the data

Incompleteness and generalisation



<https://kids.strw.leidenuniv.nl/>




AI-Driven HEP 7

What does an Input look like

An array with shape (n, m, l, k)

Domain Distribution





AI-Driven HEP 7

What does an Input look like

An array with shape (n, m, l, k)

Domain Distribution








AI-Driven HEP 7

What does an Input look like

An array with shape (n, m, l, k)

Domain Distribution







AI-Driven HEP 7

What does an Input look like

An array with shape (n, m, l, k)

Domain Uncertainty






AI-Driven HEP 7


What does an Input look like

An array with shape (n, m, l, k)

Domain Uncertainty

I.I.D





AI-Driven HEP 7


What does an Input look like


An array with shape (n, m, l, k)

Domain Uncertainty

I.I.D

$$\mathcal{N}(x|\mu, \sigma^2) = \frac{1}{(2\pi\sigma^2)^{1/2}} \exp \left\{ -\frac{1}{2\sigma^2} (x - \mu)^2 \right\}$$





AI-Driven HEP 7


What does an Input look like

An array with shape (n, m, l, k)


Domain Uncertainty

I.I.D

$$\mathcal{N}(x|\mu, \sigma^2) = \frac{1}{(2\pi\sigma^2)^{1/2}} \exp \left\{ -\frac{1}{2\sigma^2} (x - \mu)^2 \right\}$$



$$p(\mathbf{x}|\mu, \sigma^2) = \prod_{n=1}^N \mathcal{N}(x_n|\mu, \sigma^2)$$




AI-Driven HEP 7


What does an Input look like

An array with shape (n, m, l, k)

Domain Uncertainty

I.I.D ?






AI-Driven HEP 7


What does an Input look like

An array with shape (n, m, l, k)

Domain Uncertainty

Prediction (Inference)





AI-Driven HEP 7


What does an Input look like

An array with shape (n, m, l, k)


Domain Uncertainty

Prediction (Inference)

$$\hat{Y} = \hat{f}(X)$$



$$Y = f(X) + \epsilon$$



AI-Driven HEP 7


What does an Input look like

An array with shape (n, m, l, k)

Domain Uncertainty


Prediction (Inference)

$$\hat{Y} = \hat{f}(X)$$



$$Y = f(X) + \epsilon$$

$$\begin{aligned} \mathbb{E}(Y - \hat{Y})^2 &= \mathbb{E}[f(X) + \epsilon - \hat{f}(X)]^2 \\ &= \underbrace{\mathbb{E}[f(X) - \hat{f}(X)]^2}_{\text{Reducible}} + \underbrace{\mathbb{E}[\epsilon^2]}_{\text{Irreducible}} \end{aligned}$$



AI-Driven HEP 7


What does an Input look like

An array with shape (n, m, l, k)

Domain Uncertainty

Prediction (Inference)


$$\hat{Y} = \hat{f}(X)$$



$$Y = f(X) + \epsilon$$

$$\begin{aligned} E(Y - \hat{Y})^2 &= E[f(X) + \epsilon - \hat{f}(X)]^2 \\ &= \underbrace{[f(X) - \hat{f}(X)]^2}_{\text{Reducible}} + \underbrace{\text{Var}(\epsilon)}_{\text{Irreducible}} \end{aligned}$$

Show the MLE reduce to squared error loss function



AI-Driven HEP 7

What does an Input look like

An array with shape (n, m, l, k)

Bias - Variance Tradeoff

