





Introduction to


AI-Driven HEP

S. A. Fard - School of Physics (IPM)






AI-Driven HEP 9




Session 9

Neural Networks



Reference

- [1] James, Gareth, et al. "Statistical learning." An introduction to statistical learning: With applications in Python. (2023)
- [2] Hastie, Trevor. "The elements of statistical learning: data mining, inference, and prediction." (2009)
- [3] Bishop, Christopher M., Pattern recognition and machine learning. springer, (2006)
- [4] Chollet, Francois. Deep learning with Python. (2021)
- [5] Shalev-Shwartz, Shai, and Shai Ben-David. Understanding machine learning: From theory to algorithms. (2014)




AI-Driven HEP 9

Motivation

Remember the challenges of previous session

Finding the optimum value for parameters


Generalizing the approach to more complicated functions




AI-Driven HEP 9

Optimization


$$y = ax + b$$






Optimization

$$y = ax + b$$




Loss Function




Squared-Error

$$L = \sum_i (y_i - ax_i - b)^2$$




Optimization

$$y = ax + b$$




Loss Function



Squared-Error


$$L = \sum_i (y_i - ax_i - b)^2$$

More ? wait



Optimization


$$y = ax + b$$



Gradient Descent


$$a_{new} = a_{old} - \frac{\partial L(a_{old}, b_{old})}{\partial a_{old}}$$

$$b_{new} = b_{old} - \frac{\partial L(a_{old}, b_{old})}{\partial b_{old}}$$



Optimization

$$y = ax + b$$




Gradient Descent

$$a_{new} = a_{old} - \frac{\partial L(a_{old}, b_{old})}{\partial a_{old}}$$

$$b_{new} = b_{old} - \frac{\partial L(a_{old}, b_{old})}{\partial b_{old}}$$


Solve for this data points
starting with (a=1,b=1)

$x_1 = 1$	$y_1 = 1$
$x_2 = 2$	$y_2 = 2$
$x_3 = 3$	$y_3 = 2$



Optimization

$$y = ax + b$$




Gradient Descent

$$a_{new} = a_{old} - \eta \frac{\partial L(a_{old}, b_{old})}{\partial a_{old}}$$


$$b_{new} = b_{old} - \eta \frac{\partial L(a_{old}, b_{old})}{\partial b_{old}}$$

There is a need for the learning rate



Optimization

$$y = ax + b$$




$$a_{new} = a_{old} - \eta \frac{\partial L(a_{old}, b_{old})}{\partial a_{old}}$$

Stochastic Gradient Descent


$$b_{new} = b_{old} - \eta \frac{\partial L(a_{old}, b_{old})}{\partial b_{old}}$$

The whole data set or a
subset of data (mini-batch)



Optimization

$$y = ax + b$$




$$a_{new} = a_{old} - \eta \frac{\partial L(a_{old}, b_{old})}{\partial a_{old}}$$

Stochastic Gradient Descent

$$b_{new} = b_{old} - \eta \frac{\partial L(a_{old}, b_{old})}{\partial b_{old}}$$


More ?

For example see Adam Optimizer



AI-Driven HEP 9


Generalization




Generalization

Adding more features

$$y = a_1x_1 + a_2x_2 + b$$







Generalization

Adding more features
Capturing Non-linearity

$$y = A(a_1x_1 + a_2x_2 + b)$$






Generalization

Adding more features
Capturing Non-linearity

$$y = A(a_1x_1 + a_2x_2 + b)$$



Heaviside function


$$H(X) = \begin{cases} 1 & x \geq 0 \\ 0 & x < 0 \end{cases}$$

Sigmoid


$$g(z) = \frac{e^z}{1 + e^z} = \frac{1}{1 + e^{-z}}$$

Rectified Linear Unit (ReLU)

$$\sigma(x) = \max(0, x)$$



Generalization




ReLU

Simple derivative

Fast calculation


Sigmoid


convert a linear function
into probabilities between
zero and one



Generalization


Adding more features
Capturing Non-linearity
Adding more Hidden nodes







Generalization


Adding more features
Capturing Non-linearity
Adding more Hidden nodes






Neural Network







AI-Driven HEP 9

Neural Network






Neural Network




Universal approximation theorem

There exists a neural network with one hidden layer with a finite number of nodes and a non-polynomial activation function approximate any reasonable function to arbitrary precision




Neural Network




Universal approximation theorem

There exists a neural network with one hidden layer with a finite number of nodes and a non-polynomial activation function approximate any reasonable function to arbitrary precision

- Large enough network can approximate any function
- Only guarantees existence, not being able to find it
- In practice: one huge layer not great for training




Neural Network



Universal approximation theorem


There exists a neural network with one hidden layer with a finite number of nodes and a non-polynomial activation function approximate any reasonable function to arbitrary precision


- Large enough network can approximate any function
- Only guarantees existence, not being able to find it
- In practice: one huge layer not great for training



AI-Driven HEP 9

Neural Network






Softmax

For a vector of logits $z = [z_1, z_2, \dots, z_C]$

$$\text{softmax}(z_i) = \frac{e^{z_i}}{\sum_{j=1}^C e^{z_j}}$$

Logits are unnormalized scores

C is the number of output (classes)



Softmax

For a vector of logits $z = [z_1, z_2, \dots, z_C]$

$$\text{softmax}(z_i) = \frac{e^{z_i}}{\sum_{j=1}^C e^{z_j}}$$

Logits are unnormalized scores


C is the number of output (classes)

z_i can be any real number (positive or negative)

e^{z_i} makes all values positive and amplifies larger ones.

outputs are between 0 and 1, and together they sum to 1


each output i can be interpreted as the probability that the input belongs to class i



Cross-Entropy Loss

What does this function called in physics ?

$$\text{softmax}(z_i) = \frac{e^{z_i}}{\sum_{j=1}^C e^{z_j}}$$




Cross-Entropy Loss

What does this function called in physics ?

$$\text{softmax}(z_i) = \frac{e^{z_i}}{\sum_{j=1}^C e^{z_j}}$$

The Boltzmann (Gibbs) distribution

$$p_i = \frac{1}{Q} \exp\left(-\frac{\varepsilon_i}{k_B T}\right) = \frac{\exp\left(-\frac{\varepsilon_i}{k_B T}\right)}{\sum_{j=1}^M \exp\left(-\frac{\varepsilon_j}{k_B T}\right)}$$



Cross-Entropy Loss

What does this function called in physics ?

$$\text{softmax}(z_i) = \frac{e^{z_i}}{\sum_{j=1}^C e^{z_j}}$$

The Boltzmann (Gibbs) distribution

$$p_i = \frac{1}{Q} \exp\left(-\frac{\varepsilon_i}{k_B T}\right) = \frac{\exp\left(-\frac{\varepsilon_i}{k_B T}\right)}{\sum_{j=1}^M \exp\left(-\frac{\varepsilon_j}{k_B T}\right)}$$

Ex3 : Prove that provided the Boltzmann distribution the loss function is given by

$$L = -\frac{1}{N} \sum_{n=1}^N \sum_{i=1}^c y_{n,i} \log(\hat{y}_{n,i})$$